

Department of Electronic and Telecommunication Engineering  
University of Moratuwa



EN4554 - Deep Learning for Vision

**Assignment 02**

|                          |           |
|--------------------------|-----------|
| Abeygunathilaka T. L.    | - 200003P |
| De Silva W. H. P.        | - 200114G |
| Haputhanthri H. H. A. M. | - 200207U |
| Ranasingha A. S. N.      | - 200507N |

## Part 01

When broadcasting two arrays, NumPy compares their shapes, starting from the rightmost dimension to the left. Two dimensions are compatible when;

- They are equal, or
- One of them is 1

The number of dimensions need not be equal. In unequal dimension count cases, they are “stretched” to make the computation compatible.




Figure 01: [Unequal dimension count](#)

- I.  $a + b \rightarrow (256, 256, 3)$
- II.  $a - b \rightarrow (10, 9, 6, 7)$
- III.  $a * b \rightarrow \text{Error}$
- IV.  $a / b \rightarrow \text{Error}$
- V.  $a - b \rightarrow (5, 3, 2)$
- VI.  $a - b.\text{mean}() \rightarrow (128, 128, 3)$
- VII.  $a - b.\text{mean}(\text{axis}=(1, 2)) \rightarrow \text{Error}$
- VIII.  $a - b.\text{mean}(\text{axis}=(1, 2), \text{keepdims}=\text{True}) \rightarrow (3, 256, 256)$
- IX.  $\text{np.matmul}(a, b) \rightarrow (6, 5, 4)$

## Part 02

The formula for the squared Euclidean distance between two vectors  $x_i$  and  $y_j$  is given by;

$$\|x_i - y_j\|^2 = \sum_{k=1}^d x_{i,k}^2 + \sum_{k=1}^d y_{j,k}^2 - 2 \sum_{k=1}^d x_{i,k} y_{j,k}$$

```
import numpy as np

def pairwise_squared_euclidean(X, Y):
    # Compute the squared norms for each row in X and Y using broadcasting
    X_norm = np.sum(X**2, axis=1) # Shape (m,)
    Y_norm = np.sum(Y**2, axis=1) # Shape (n,)

    # Compute the pairwise squared Euclidean distances using broadcasting and einsum
    Z = X_norm[:, np.newaxis] + Y_norm - 2 * np.einsum('ij,kj->ik', X, Y)

    # Ensure non-negative distances (can be slightly negative due to floating-point precision)
    Z = np.maximum(Z, 0)

    return Z
```

Results of the function for arbitrary x,y matrices,

```
x = np.array([[1,2,3],[2,3,4],[4,5,6]])
y = np.array([[2,4,5],[7,6,4],[1,1,1],[4,3,2]])
print(pairwise_squared_euclidean3(x, y))
```

✓ 0.0s

```
[[ 9 53  5 11]
 [ 2 34 14  8]
 [ 6 14 50 20]]
```

## **Part 03**

**(a)**

Find optimal k using inbuilt functions


```
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, train_embeddings, train_labels, cv=3, scoring='accuracy')
    cv_scores.append(scores.mean())
optimal_k = k_values[np.argmax(cv_scores)]
print(f"Optimal value of k: {optimal_k}")
```

Find optimal k using custom functions and find accuracy

```
def pairwise_squared_euclidean_distance(X, Y):
    X_sq_norms = np.sum(X ** 2, axis=1).reshape(-1, 1)
    Y_sq_norms = np.sum(Y ** 2, axis=1).reshape(1, -1)
    cross_term = np.dot(X, Y.T)
    dists = X_sq_norms + Y_sq_norms - 2 * cross_term
    return dists
```

```
def knn_predict(X_train, train_labels, X_test, k=5):
    dists = pairwise_squared_euclidean_distance(X_test, X_train)
    num_test_samples = X_test.shape[0]
    y_pred = np.zeros(num_test_samples, dtype=int)
    for i in range(num_test_samples):
        nearest_neighbors = np.argsort(dists[i])[:k]
        nearest_labels = train_labels[nearest_neighbors]
        unique, counts = np.unique(nearest_labels, return_counts=True)
        y_pred[i] = unique[np.argmax(counts)]
    return y_pred
```

Performed 3-fold cross-validation to find the optimum value of k. The optimum k value is 3.



The accuracy on the test set is 84.16% when k = 3

**(b)** Linear classification

```
# Load train and test embeddings
train_embeddings = np.load('/content/drive/MyDrive/Colab Notebooks/embeddings/train_embeddings.npy')
train_labels = np.load('/content/drive/MyDrive/Colab Notebooks/embeddings/train_labels.npy')
test_embeddings = np.load('/content/drive/MyDrive/Colab Notebooks/embeddings/test_embeddings.npy')
test_labels = np.load('/content/drive/MyDrive/Colab Notebooks/embeddings/test_labels.npy')

# Train a Logistic Regression model (equivalent to the last fully connected layer)
log_reg = LogisticRegression(solver='lbfgs', max_iter=1000)

# Fit the logistic regression model on the training embeddings
log_reg.fit(train_embeddings, train_labels)

# Predict on the test embeddings
test_predictions = log_reg.predict(test_embeddings)

# Calculate accuracy
accuracy = accuracy_score(test_labels, test_predictions)
print(f'Accuracy on the test set (Logistic Regression): {accuracy * 100:.2f}%')
```

Accuracy on the test set (Logistic Regression): 92.41%

(c) Part C is written in PyTorch separately in Kaggle

(i) Initiating the model

```
class ImageModel(nn.Module):
    def __init__(self):
        super(ImageModel, self).__init__()
```

(ii) Forward function

```
def forward(self, img):
    feats = self.features(img)
    x = self.flatten(feats)
    x = self.dense(x)

    out = F.softmax(x, dim=1)
    return out
```

(iii) Model training

```
for (img, label) in train_loader:
    img, label = img.to(device), label.to(device)
    one_hot_labels = one_hot_encode(label)

    optimizer.zero_grad()

    pred = model(img)
    loss = criterion(pred, one_hot_labels)
    loss.backward()

    optimizer.step()
    batch_loss += loss.item()

    with torch.no_grad():
        y_pred = torch.argmax(pred, dim=1)
        total += label.size(0) # Update total count
        correct += (y_pred == label).sum().item()

train_loss.append(batch_loss/len(train_loader))
train_acc.append(correct/total)
```

(iv) Test accuracy evaluation

```
with torch.no_grad():
    for (img, label) in test_loader:
        img, label = img.to(device), label.to(device)
        one_hot_labels = one_hot_encode(label)

        pred = model(img)
        loss = criterion(pred, one_hot_labels)
        batch_loss += loss.item()

        y_pred = torch.argmax(pred, dim=1)
        total += label.size(0) # Update total count
        correct += (y_pred == label).sum().item()

test_loss.append(batch_loss/len(test_loader))
test_acc.append(correct/total)
```

(v) Training accuracy logs

|          |  |     |   |
|----------|--|-----|---|
| epoch 0  |  | 100 | train_loss= 4.3217 , train_acc= 0.3190, test_loss: 4.5038, test_acc= 0.1210 |
| epoch 1  |  | 100 | train_loss= 4.2841 , train_acc= 0.3491, test_loss: 4.3238, test_acc= 0.3128 |
| epoch 2  |  | 100 | train_loss= 4.2683 , train_acc= 0.3645, test_loss: 4.2813, test_acc= 0.3546 |
| epoch 3  |  | 100 | train_loss= 4.2653 , train_acc= 0.3677, test_loss: 4.2614, test_acc= 0.3771 |
| epoch 4  |  | 100 | train_loss= 4.2584 , train_acc= 0.3731, test_loss: 4.3066, test_acc= 0.3332 |
| epoch 5  |  | 100 | train_loss= 4.2616 , train_acc= 0.3724, test_loss: 4.3070, test_acc= 0.3294 |
| epoch 6  |  | 100 | train_loss= 4.2397 , train_acc= 0.3928, test_loss: 4.2693, test_acc= 0.3674 |
| epoch 7  |  | 100 | train_loss= 4.2304 , train_acc= 0.4023, test_loss: 4.2426, test_acc= 0.3944 |
| epoch 8  |  | 100 | train_loss= 4.2261 , train_acc= 0.4058, test_loss: 4.2434, test_acc= 0.3930 |
| epoch 9  |  | 100 | train_loss= 4.2308 , train_acc= 0.4030, test_loss: 4.2600, test_acc= 0.3775 |
| epoch 10 |  | 100 | train_loss= 4.2233 , train_acc= 0.4089, test_loss: 4.2320, test_acc= 0.4044 |
| epoch 11 |  | 100 | train_loss= 4.2166 , train_acc= 0.4167, test_loss: 4.2232, test_acc= 0.4117 |
| epoch 12 |  | 100 | train_loss= 4.2121 , train_acc= 0.4220, test_loss: 4.2154, test_acc= 0.4207 |
| epoch 13 |  | 100 | train_loss= 4.2018 , train_acc= 0.4312, test_loss: 4.2304, test_acc= 0.4044 |
| epoch 14 |  | 100 | train_loss= 4.1997 , train_acc= 0.4329, test_loss: 4.2327, test_acc= 0.4055 |
| epoch 15 |  | 100 | train_loss= 4.1887 , train_acc= 0.4440, test_loss: 4.1905, test_acc= 0.4438 |
| epoch 16 |  | 100 | train_loss= 4.1759 , train_acc= 0.4573, test_loss: 4.1914, test_acc= 0.4431 |
| epoch 17 |  | 100 | train_loss= 4.1706 , train_acc= 0.4623, test_loss: 4.2019, test_acc= 0.4293 |
| epoch 18 |  | 100 | train_loss= 4.1721 , train_acc= 0.4632, test_loss: 4.1867, test_acc= 0.4487 |
| epoch 19 |  | 100 | train_loss= 4.1658 , train_acc= 0.4677, test_loss: 4.2015, test_acc= 0.4355 |
| epoch 20 |  | 100 | train_loss= 4.1611 , train_acc= 0.4699, test_loss: 4.1924, test_acc= 0.4380 |